

Raisin Classification



by: Noura El-Mamlouk

```
In [113...  ##Importing Libraries##
import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from sklearn.utils import shuffle
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn import svm
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import make_scorer, accuracy_score
from sklearn import metrics
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import classification_report
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB
```

1)Data Preprocessing and EDA

a) Data Loading

```
In [114... raisin = pd.read_csv("Raisin_Dataset (1).csv")
raisin.head()
```

Out[114]:

	Area	MajorAxisLength	MinorAxisLength	Eccentricity	ConvexArea	Extent	Perimeter	C
0	87524	442.246011	253.291155	0.819738	90546	0.758651	1184.040	Kecir
1	75166	406.690687	243.032436	0.801805	78789	0.684130	1121.786	Kecir
2	90856	442.267048	266.328318	0.798354	93717	0.637613	1208.575	Kecir
3	45928	286.540559	208.760042	0.684989	47336	0.699599	844.162	Kecir
4	79408	352.190770	290.827533	0.564011	81463	0.792772	1073.251	Kecir

In [115... raisin.shape

Out[115]: (900, 8)

b) Data Cleaning

In [116... duplicates = raisin.duplicated()
print("Number of duplicates:", duplicates.sum())

Number of duplicates: 0

In [117... missing_values = raisin.isnull().sum()
print("Missing values:\n", missing_values)

Missing values:

```
Area          0
MajorAxisLength  0
MinorAxisLength  0
Eccentricity    0
ConvexArea      0
Extent          0
Perimeter       0
Class           0
dtype: int64
```

In [118... import matplotlib.pyplot as plt

```
# Define columns of interest
columns = ['Area', 'MajorAxisLength', 'MinorAxisLength', 'Eccentricity', 'ConvexArea']

# Calculate the IQR for each column
Q1 = raisin[columns].quantile(0.25)
Q3 = raisin[columns].quantile(0.75)
IQR = Q3 - Q1

# Define outlier threshold
outlier_threshold = 1.5

# Identify outliers in each column
outliers = ((raisin[columns] < (Q1 - outlier_threshold * IQR)) | (raisin[columns] > (Q3 + outlier_threshold * IQR)))

# Count the number of outliers in each column before removing
num_outliers_before = outliers.sum()

# Filter the DataFrame to remove outliers
raisin_filtered = raisin[~outliers.any(axis=1)]

# Calculate the IQR for each column after removing outliers
Q1_filtered = raisin_filtered[columns].quantile(0.25)
```

```

Q3_filtered = raisin_filtered[columns].quantile(0.75)
IQR_filtered = Q3_filtered - Q1_filtered

# Identify outliers in each column after removing
outliers_filtered = ((raisin_filtered[columns] < (Q1_filtered - outlier_threshold)

# Count the number of outliers in each column after removing
num_outliers_after = outliers_filtered.sum()


# Create a boxplot for each column after removing outliers
fig, axs = plt.subplots(1, len(columns), figsize=(16, 4))

# Iterate over each column and plot the boxplot
for i, column in enumerate(columns):
    axs[i].boxplot(raisin_filtered[column])
    axs[i].set_title(column)
    axs[i].set_xlabel('Column')
    axs[i].set_ylabel('Value')
    axs[i].text(0.85, 0.85, f'Outliers: {num_outliers_after[column]}', transform=axs[i].transData)

plt.tight_layout()
plt.show()

# Output the number of outliers before and after removing
print("Number of Outliers Before Removing:", num_outliers_before)
print("Number of Outliers After Removing:", num_outliers_after)
raisin = raisin_filtered
raisin.shape

```



```

Number of Outliers Before Removing: Area          41
MajorAxisLength      17
MinorAxisLength      26
Eccentricity         43
ConvexArea           42
Extent               21
Perimeter            17
dtype: int64
Number of Outliers After Removing: Area          15
MajorAxisLength       2
MinorAxisLength       2
Eccentricity          4
ConvexArea            10
Extent                6
Perimeter             1
dtype: int64
(795, 8)

```

Out[118]:

```

In [119... # Check for inconsistent values
for column in raisin.columns:
    unique_values = raisin[column].unique()

# Replace inconsistent values with the mean
raisin_cleaned = raisin.copy()

```

```

for column in raisin_cleaned.columns:
    if column != 'Class': # Exclude the 'Class' column from modification
        mean_value = raisin_cleaned[column].mean()
        raisin_cleaned[column].replace(-1, mean_value, inplace=True)

# Verify the result
print("Dataset after handling inconsistent values:")
print(raisin_cleaned.head())
raisin = raisin_cleaned

```

Dataset after handling inconsistent values:

	Area	MajorAxisLength	MinorAxisLength	Eccentricity	ConvexArea	\
0	87524	442.246011	253.291155	0.819738	90546	
1	75166	406.690687	243.032436	0.801805	78789	
2	90856	442.267048	266.328318	0.798354	93717	
3	45928	286.540559	208.760042	0.684989	47336	
5	49242	318.125407	200.122120	0.777351	51368	

	Extent	Perimeter	Class
0	0.758651	1184.040	Kecimen
1	0.684130	1121.786	Kecimen
2	0.637613	1208.575	Kecimen
3	0.699599	844.162	Kecimen
5	0.658456	881.836	Kecimen

c) Data Transformation

```

In [120... from sklearn.preprocessing import LabelEncoder
# Initialize the Label encoder
label_encoder = LabelEncoder()
# Apply Label encoding to the 'Class' column
raisin['Class'] = label_encoder.fit_transform(raisin['Class'])

```

```

In [121... raisin.head()

```

```

Out[121]:

```

	Area	MajorAxisLength	MinorAxisLength	Eccentricity	ConvexArea	Extent	Perimeter	Class
0	87524	442.246011	253.291155	0.819738	90546	0.758651	1184.040	1
1	75166	406.690687	243.032436	0.801805	78789	0.684130	1121.786	1
2	90856	442.267048	266.328318	0.798354	93717	0.637613	1208.575	1
3	45928	286.540559	208.760042	0.684989	47336	0.699599	844.162	1
5	49242	318.125407	200.122120	0.777351	51368	0.658456	881.836	1

```

In [122... # Summary statistics
print("Summary Statistics:")
print(raisin.describe())

```

Summary Statistics:

	Area	MajorAxisLength	MinorAxisLength	Eccentricity	\
count	795.000000	795.000000	795.000000	795.000000	
mean	82685.197484	419.662193	247.605251	0.790847	
std	31257.813387	95.470884	42.239791	0.066809	
min	25387.000000	225.629541	143.710872	0.596359	
25%	58949.000000	347.249548	216.722981	0.749021	
50%	77105.000000	403.909415	244.803889	0.800063	
75%	100742.000000	477.704334	275.238330	0.840699	
max	171749.000000	696.149046	369.286454	0.928094	

	ConvexArea	Extent	Perimeter	Class
count	795.000000	795.000000	795.000000	795.000000
mean	85710.872956	0.700641	1134.306253	0.513208
std	32169.943577	0.044270	224.023387	0.500140
min	26139.000000	0.576457	619.074000	0.000000
25%	61190.000000	0.671604	964.835500	0.000000
50%	80094.000000	0.705007	1106.033000	1.000000
75%	104361.500000	0.732440	1269.373000	1.000000
max	177170.000000	0.830632	1755.968000	1.000000

In [123...

```
# Define the columns to normalize
columns_to_normalize = ['Area', 'MajorAxisLength', 'MinorAxisLength', 'Eccentricity']

# Initialize the scaler
scaler = MinMaxScaler()

# Perform min-max scaling on the selected columns
raisin_scaled = raisin.copy()
raisin_scaled[columns_to_normalize] = scaler.fit_transform(raisin[columns_to_normalize])

# Print the normalized dataset
print("Normalized Dataset:")
print(raisin_scaled.head())
raisin=raisin_scaled
```

Normalized Dataset:

	Area	MajorAxisLength	MinorAxisLength	Eccentricity	ConvexArea	\
0	0.424543	0.460377	0.485781	0.673368	0.426449	
1	0.340109	0.384811	0.440303	0.619309	0.348604	
2	0.447309	0.460422	0.543576	0.608905	0.447445	
3	0.140344	0.129455	0.288370	0.267173	0.140349	
5	0.162986	0.196582	0.250077	0.545594	0.167045	

	Extent	Perimeter	Class
0	0.758651	1184.040	1
1	0.684130	1121.786	1
2	0.637613	1208.575	1
3	0.699599	844.162	1
5	0.658456	881.836	1

d) Data Reduction

In [124...

```
# Assuming 'Perimeter' and 'Extent' are irrelevant attributes
raisin.drop(['Perimeter', 'Extent'], axis=1, inplace=True)
raisin.head()
```

Out[124]:

	Area	MajorAxisLength	MinorAxisLength	Eccentricity	ConvexArea	Class
--	------	-----------------	-----------------	--------------	------------	-------

0	0.424543	0.460377	0.485781	0.673368	0.426449	1
1	0.340109	0.384811	0.440303	0.619309	0.348604	1
2	0.447309	0.460422	0.543576	0.608905	0.447445	1
3	0.140344	0.129455	0.288370	0.267173	0.140349	1
5	0.162986	0.196582	0.250077	0.545594	0.167045	1

In [125... raisin.shape

Out[125]: (795, 6)

In [126...
Summary statistics
print("Summary Statistics:")
print(raisin.describe())

Summary Statistics:

	Area	MajorAxisLength	MinorAxisLength	Eccentricity	ConvexArea	\
count	795.000000	795.000000	795.000000	795.000000	795.000000	
mean	0.391483	0.412380	0.460575	0.586276	0.394435	
std	0.213565	0.202905	0.187253	0.201392	0.213002	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.229308	0.258480	0.323670	0.460194	0.232078	
50%	0.353357	0.378900	0.448156	0.614057	0.357245	
75%	0.514854	0.535737	0.583075	0.736553	0.517923	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	Class
count	795.000000
mean	0.513208
std	0.500140
min	0.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	1.000000

In [127...
Bar plot
sns.countplot(data=raisin, x='Class')
plt.title('Class Distribution')
plt.xlabel('Class')
plt.ylabel('Count')
plt.show()

Histogram
sns.histplot(data=raisin, x='Area', bins=20)
plt.title('Distribution of Area')
plt.xlabel('Area')
plt.ylabel('Count')
plt.show()



Scatter plot
sns.scatterplot(data=raisin, x='MajorAxisLength', y='MinorAxisLength', hue='Class')
plt.title('Major Axis Length vs Minor Axis Length')
plt.xlabel('Major Axis Length')
plt.ylabel('Minor Axis Length')
plt.show()



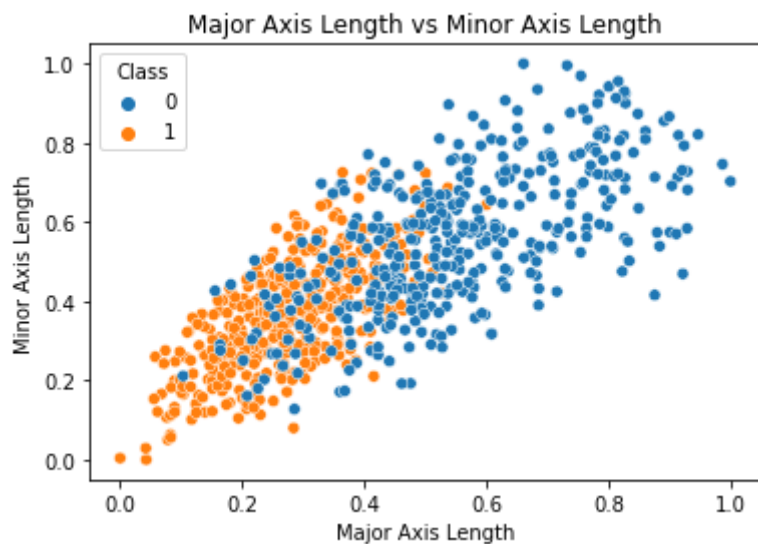
Box plot
sns.boxplot(data=raisin, x='Class', y='Area')

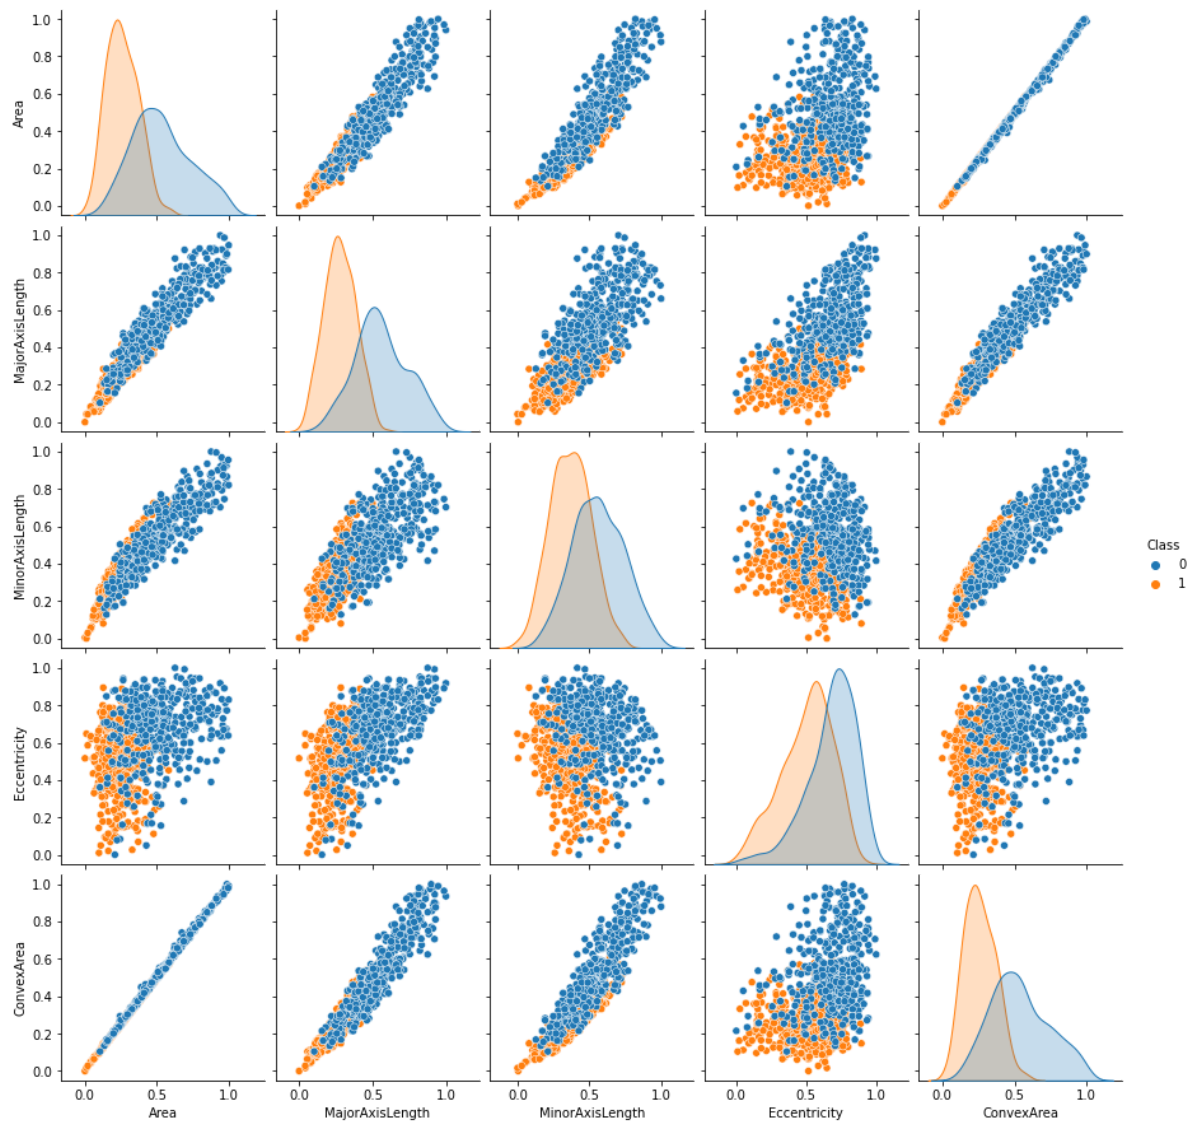
```
plt.title('Area Distribution by Class')
plt.xlabel('Class')
plt.ylabel('Area')
plt.show()

# Correlation matrix heatmap
corr_matrix = raisin.corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()

# Pairwise scatter plot matrix
sns.pairplot(data=raisin, vars=['Area', 'MajorAxisLength', 'MinorAxisLength', 'Ecc'])
plt.show()
```







2) Supervised Machine Learning

```
In [128... from sklearn.model_selection import GridSearchCV
def preprocess_and_compare(raisin, num_samples):
    """
    Preprocess the data and compare the performance of KNN, Decision Tree, and Naïve
    Parameters:
    raisin (DataFrame): The original dataframe containing the raisin data.
    num_samples (int): The number of samples to be taken from the "Besni" class.
    Returns:
    None
    """

    # Randomly select the specified number of samples from the "Besni" class
    sample_of_besni = besni.sample(n=num_samples, random_state=42)

    # Concatenate the sampled "Besni" class with the "Kecimen" class to create the
    new_df = pd.concat([sample_of_besni, kecimen], axis=0)

    # Split the data into features and labels
    X = new_df.iloc[:, :-1]
    y = new_df.iloc[:, -1]

    # Split the data into training and test sets
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, shuffle=True)

compare_classifiers(X_train, X_test, y_train, y_test)
def compare_classifiers(X_train, X_test, y_train, y_test):
    """
    Compare the performance of KNN, Decision Tree, and Naïve Bayes classifiers with
    Parameters:
    X_train (ndarray): The training input samples.
    X_test (ndarray): The testing input samples.
    y_train (ndarray): The training target samples.
    y_test (ndarray): The testing target samples.

    Returns:
    None
    """

    classifiers = {
        'KNN': {
            'model': KNeighborsClassifier(),
            'params': {'n_neighbors': [3, 5, 7]}
        },
        'Decision Tree': {
            'model': DecisionTreeClassifier(),
            'params': {'max_depth': [None, 3, 5]}
        },
        'Naïve Bayes': {
            'model': GaussianNB(),
            'params': {}
        }
    }

    results = []
    for name, classifier in classifiers.items():
        # Perform hyperparameter tuning using GridSearchCV
        model = GridSearchCV(classifier['model'], classifier['params'], cv=5)
        model.fit(X_train, y_train)

        # Make predictions on the test set using the best estimator
        y_pred = model.best_estimator_.predict(X_test)

        # Calculate performance metrics
        accuracy = metrics.accuracy_score(y_test, y_pred)
        precision = metrics.precision_score(y_test, y_pred)
        recall = metrics.recall_score(y_test, y_pred)
        f1 = metrics.f1_score(y_test, y_pred)
        results.append({'classifier': name, 'accuracy': accuracy})

        # Print the metrics
        print(f"{name} Accuracy: {accuracy:.3f}")
        print(f"{name} Precision: {precision:.3f}")
        print(f"{name} Recall: {recall:.3f}")
        print(f"{name} F1: {f1:.3f}")

        # Create the confusion matrix
        cm = metrics.confusion_matrix(y_test, y_pred)

        # Plot the confusion matrix
        plt.figure()
        sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square=True, cmap='Blues')
        plt.ylabel('Actual label')
        plt.xlabel('Predicted label')
        plt.title(f'{name} Confusion Matrix')
        plt.show()

```


```
# Visualize the accuracy of each classifier
plt.figure(figsize=(10, 6))
sns.barplot(x="classifier", y="accuracy", data=pd.DataFrame(results))
plt.ylim(0, 1)
plt.ylabel("Accuracy")
plt.title("Classifier Accuracy Comparison")
plt.show()

highest_accuracy = max(results, key=lambda x: x['accuracy'])
print(f"The classifier with the highest accuracy is {highest_accuracy['classifier']}")


lowest_accuracy = min(results, key=lambda x: x['accuracy'])
print(f"The classifier with the lowest accuracy is {lowest_accuracy['classifier']}")
```

In [129... preprocess_and_compare(raisin, 100)

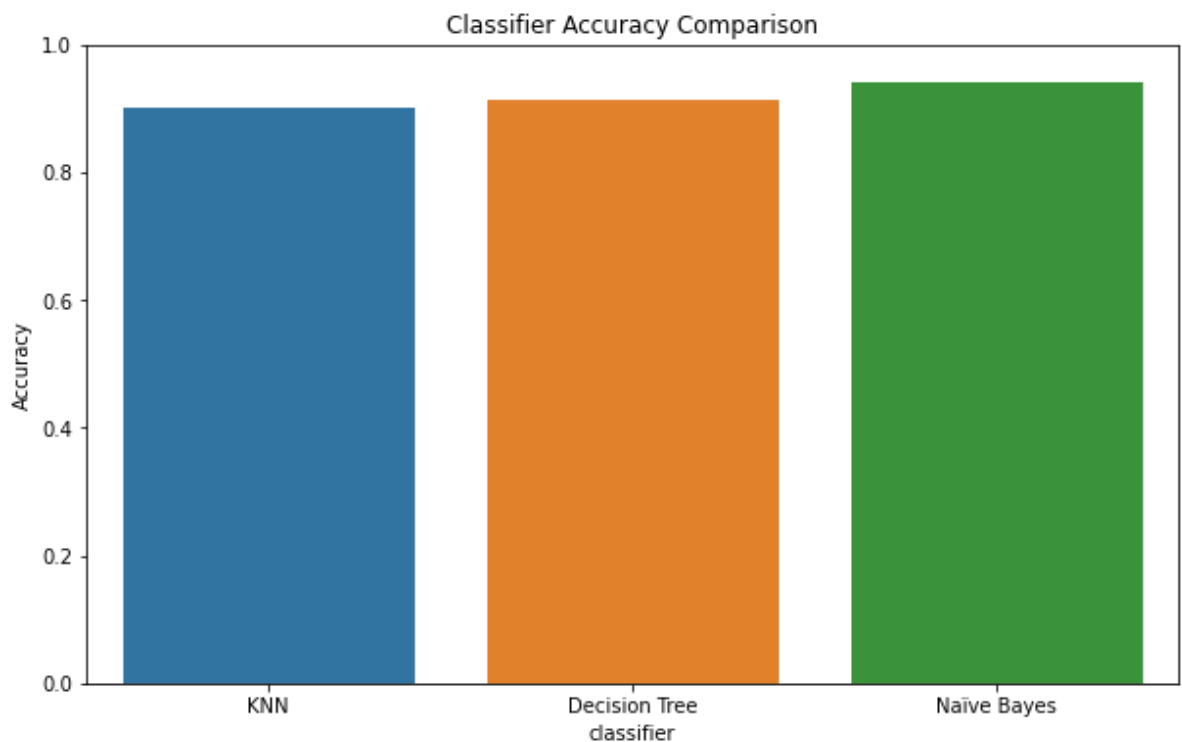

KNN Accuracy: 0.902
 KNN Precision: 0.885
 KNN Recall: 1.000
 KNN F1: 0.939



Decision Tree Accuracy: 0.912
 Decision Tree Precision: 0.905
 Decision Tree Recall: 0.987
 Decision Tree F1: 0.944



Naïve Bayes Accuracy: 0.941
 Naïve Bayes Precision: 0.928
 Naïve Bayes Recall: 1.000
 Naïve Bayes F1: 0.963




The classifier with the highest accuracy is Naïve Bayes with an accuracy of 0.941
The classifier with the lowest accuracy is KNN with an accuracy of 0.902


In [130...

```
preprocess_and_compare(raisin, 200)
```


```
KNN Accuracy: 0.844
KNN Precision: 0.839
KNN Recall: 0.936
KNN F1: 0.885
```




Decision Tree Accuracy: 0.869
Decision Tree Precision: 0.878
Decision Tree Recall: 0.923
Decision Tree F1: 0.900



Naïve Bayes Accuracy: 0.869
Naïve Bayes Precision: 0.860
Naïve Bayes Recall: 0.949
Naïve Bayes F1: 0.902





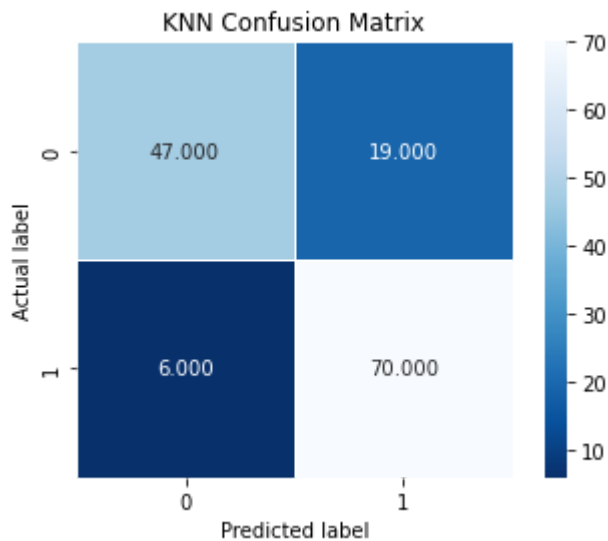
The classifier with the highest accuracy is Decision Tree with an accuracy of 0.869

The classifier with the lowest accuracy is KNN with an accuracy of 0.844


In [131...

```
preprocess_and_compare(raisin, 300)
```



KNN Accuracy: 0.824
 KNN Precision: 0.787
 KNN Recall: 0.921
 KNN F1: 0.848



Decision Tree Accuracy: 0.866
 Decision Tree Precision: 0.870
 Decision Tree Recall: 0.882
 Decision Tree F1: 0.876



Naïve Bayes Accuracy: 0.852
Naïve Bayes Precision: 0.831
Naïve Bayes Recall: 0.908
Naïve Bayes F1: 0.868



The classifier with the highest accuracy is Decision Tree with an accuracy of 0.866

The classifier with the lowest accuracy is KNN with an accuracy of 0.824